

```

/* =====
 *
 * CS 566 - Assignment 03
 * Camillo Lugaresi, Cosmin Stroe
 *
 * This code implements the DNS (Dekel, Nassimi, Sahni) matrix multiplication
 * algorithm on a cluster using MPI.
 *
 * The DNS algorithm uses a 3D Mesh to partition the intermediate data
 * of the matrix multiplication problem.
 *
 * ===== */

#include "mpi.h"
#include "common.h"
#include "cannon.h" // for struct problem
#include <stdio.h>
#include <stdlib.h>
#include <unistd.h>
#include <strings.h>
#include <string.h>

#define iDIM 0
#define jDIM 1
#define kDIM 2

mesh_info {
    MPI_Comm mesh3d;
    MPI_Comm mesh_ik, ring_j;
    MPI_Comm mesh_jk, ring_i;
    MPI_Comm mesh_ij, ring_k;

    int myrank;
    int coords[3]; // my coordinates in the 3D mesh

    int numprocs; // the number of processors in our cluster

    matrix mtx; // the matrix we are multiplying (n x n)

    int q; // the number of processors in each dimension of the 3D-Mesh
           // each processor will receive two blocks of (n/q)*(n/q) elements

    int n; // the size of the matrix dimension

    int k; // the power
};

void create_topology(          mesh_info *info , int argc, char *argv[] ) {
    // MPI Initialization
    MPI_Init(&argc, &argv);
    MPI_Comm_size(MPI_COMM_WORLD, &info->numprocs);

    ( info->q = 1; (info->q+1)*(info->q+1)*(info->q+1) <= info->numprocs; info->q++ );

    // Create the Topology which we will be using.
    int *dims = malloc(      (int) * 3 ); // 3 dimensions
    int *periods = malloc(    (int) * 3 ); // wraparound
    dims[iDIM] = dims[jDIM] = dims[kDIM] = info->q;
    periods[iDIM] = periods[jDIM] = periods[kDIM] = 1;

    MPI_Cart_create(MPI_COMM_WORLD, 3, dims, periods, 0, &info->mesh3d);

    MPI_Comm_rank( info->mesh3d, &info->myrank);
    MPI_Cart_coords( info->mesh3d, info->myrank, 3, info->coords);

```

```

// the i-k plane
dims[iDIM] = dims[kDIM] = 1;
dims[jDIM] = 0;

MPI_Cart_sub( info->mesh3d, dims, &info->mesh_ik );

// the j ring
dims[iDIM] = dims[kDIM] = 0;
dims[jDIM] = 1;

MPI_Cart_sub( info->mesh3d, dims, &info->ring_j);

// the j-k plane
dims[jDIM] = dims[kDIM] = 1;
dims[iDIM] = 0;

MPI_Cart_sub( info->mesh3d, dims, &info->mesh_jk);

//the i ring
dims[jDIM] = dims[kDIM] = 0;
dims[iDIM] = 1;

MPI_Cart_sub( info->mesh3d, dims, &info->ring_i);

// the i-j plane
dims[iDIM] = dims[jDIM] = 1;
dims[kDIM] = 0;

MPI_Cart_sub( info->mesh3d, dims, &info->mesh_ij);

// the k rings
dims[iDIM] = dims[jDIM] = 0;
dims[kDIM] = 1;

MPI_Cart_sub( info->mesh3d, dims, &info->ring_k);
}

/*
Distribute the left or right hand matrix.
*/
void distribute_matrix(      mesh_info *info,      matrix *A,      matrix *Asub, int ringdim, int
bycolumn) {

    int *Ablocks;
    int blksize = info->n / info->q;
    alloc_matrix(Asub, blksize);

    MPI_Comm mesh, ring;
    (ringdim == jDIM) {
        mesh = info->mesh_ik;
        ring = info->ring_j;
    }
    {
        mesh = info->mesh_jk;
        ring = info->ring_i;
    }

    ( info->myrank == 0 ) {
        Ablocks = malloc(      (*A->data) * A->n * A->n );
        matrix_to_blocks(A, Ablocks, blksize, bycolumn);
    }

    ( info->coords[ringdim] == 0 ) {
        MPI_Scatter( Ablocks, blksize*blksize, MPI_INT, Asub->data, blksize*blksize, MPI_INT, 0, mesh );
    }

    ( info->myrank == 0 ) free (Ablocks);
}

```

```

    MPI_Bcast( Asub->data, blksz*blksz, MPI_INT, 0, ring);
}

void gather_C(      mesh_info *info,      matrix *Cred,      matrix *C) {
    int *Cblocks;
    int blksz = info->n / info->q;

    ( info->myrank == 0 ) Cblocks = calloc(      (*C->data), C->n * C->n );

    MPI_Gather( Cred->data, blksz*blksz, MPI_INT, Cblocks, blksz*blksz, MPI_INT, 0, info->mesh_ij);
    ( info->myrank == 0 ) {
        blocks_to_matrix(C, Cblocks, blksz, 0);
        free(Cblocks);
    }
}

/*
This is the DNS algorithm implementation.
The multiplied matrix C will only be returned on the root (rank = 0) node.
*/
void dns_multiply(      mesh_info *info,      matrix *C,      matrix *A,      matrix *B, int
strassen)
{
    matrix Asub, Bsub, Csub, Cred;

    // distribute A along the i-k plane and then in the j direction
    distribute_matrix(info, A, &Asub, jDIM, 0);

    // distribute B along the k-j plane and then in the i direction
    distribute_matrix(info, B, &Bsub, iDIM, 1);

    alloc_matrix(&Csub, Asub.n);

    // do the serial matrix multiplication
    bzero(Csub.data,      (*Csub.data)*Csub.n*Csub.n);
    ( strassen ) {
        strassen_matrix2_mult(&Csub, &Asub, &Bsub);
    }
    {
        naive_matrix_mult_add(&Csub, &Asub, &Bsub);
    }

    // reduce along k dimension to the i-j plane
    ( info->coords[kDIM] == 0 ) alloc_matrix(&Cred, Csub.n);
    MPI_Reduce( Csub.data, Cred.data, Csub.n*Csub.n, MPI_INT, MPI_SUM, 0, info->ring_k);

    // gather on the i-j plane to the root node.
    ( info->coords[kDIM] == 0 ) gather_C(info, &Cred, C);
}

int main( int argc, char *argv[] ) {
    input_params m_in_s, *m_in = &m_in_s;
    mesh_info info;

    create_topology(&info, argc, argv);

    int result = parse_args(argc, argv, &info.n, &info.k, m_in);
    (result != 0) {
        MPI_Finalize();
        exit(result);
    }
}

```

```

    double start_time = MPI_Wtime();

    // error checking
    ( info.q*info.q*info.q != info.numprocs ) {
        fprintf(stderr, "ERROR: The number of processors must be a perfect cube (not %d).\n",
info.numprocs);
        MPI_Finalize();
        1;
    }

    // required for calling the LU decomposition functions.
    problem pinfo;
    pinfo.n = info.n;
    pinfo.p = info.numprocs;
    pinfo.k = info.k;
    pinfo.rank = info.myrank;

    alloc_matrix(&pinfo.X, pinfo.n);
    fill_matrix(&pinfo.X, m_in);

    double load_time = MPI_Wtime();

    matrix B, C;
    ( info.myrank == 0 ) {
        alloc_matrix(&C, info.n);
        alloc_matrix(&B, info.n);
        memcpy(B.data, pinfo.X.data,          (*pinfo.X.data)*pinfo.X.n*pinfo.X.n);
    }

    int i;
    ( i = 1; i < info.k; i++ ) {
        dns_multiply( &info, &C, &pinfo.X, &B, 0);
        ( info.myrank == 0 ) {
            memcpy(B.data, C.data,          (*pinfo.X.data)*pinfo.X.n*pinfo.X.n);
        }
    }

    pinfo.Xpow.data = C.data;
    pinfo.Xpow.n = C.n;

    double dns_time = MPI_Wtime();

    number_type determinant = luld_determinant(&pinfo);

    double lu_time = MPI_Wtime();

    ( info.myrank == 0 ) {
        ( m_in->print ) {
            printf("X:\n");
            print_matrix(&pinfo.X);
            printf("X^%d:\n", pinfo.k);
            print_matrix(&pinfo.Xpow);
        }
        printf("determinant: %f\n", determinant);

        double elapsed_time = MPI_Wtime() - start_time;
        printf("data loading time: %f\n", load_time - start_time);
        printf("matrix multiplication time: %f\n", dns_time - load_time);
        printf("LU time: %f\n", lu_time - dns_time);
        printf("total time: %f\n", elapsed_time);
    }

    MPI_Barrier(info.mesh3d); // avoid processors that end early from killing the rest.

    MPI_Finalize();

```

θ ;

}